

# Support Vector Machines

Yifei Sun

## Contents

<b>Classification</b>	<b>2</b>
Using <code>e1071</code> . . . . .	2
Using <code>kernlab</code> - another commonly used library in R for SVM . . . . .	9
Using <code>caret</code> . . . . .	10
Understanding your models with <code>DALEX</code> . . . . .	14
<b>Regression</b>	<b>18</b>
<code>eps-regression</code> . . . . .	18

```
library(mlbench)
library(ISLR)

library(caret) #only allows implementation from kernel lab lib
library(e1071)
library(kernlab)

library(DALEX) #generic, can be applied to lots of models
```

## Classification

We use the Pima Indians Diabetes Database for illustration. The data contain 768 observations and 9 variables. The outcome is a binary variable **diabetes**.

```
data(PimaIndiansDiabetes)
dat <- PimaIndiansDiabetes
dat$diabetes <- factor(dat$diabetes, c("pos", "neg"))

set.seed(2021)
rowTrain <- createDataPartition(y = dat$diabetes,
                                p = 0.75,
                                list = FALSE)
```

### Using e1071

Check <https://cran.r-project.org/web/packages/e1071/vignettes/svmdoc.pdf> for more details.


### Linear boundary

Most real data sets will not be fully separable by a linear boundary. Support vector classifiers with a tuning parameter **cost**, which quantifies the penalty associated with having an observation on the wrong side of the classification boundary, can be used to build a linear boundary.

```
set.seed(1)
linear.tune <- tune.svm(diabetes ~ . ,
                       data = dat[rowTrain,],
                       kernel = "linear", #give us linear decision boundary.
                       cost = exp(seq(-5,2,len=50)), #tuning parameter
                       scale = TRUE) #recommend scaling for svm
# another argument: tunecontrol = tune.control (cross = 10)

plot(linear.tune)
```

## Performance of 'svm'



```
# summary(linear.tune)
linear.tune$best.parameters #smallest cost => model is less flexible

##          cost
## 13 0.03741385

best.linear <- linear.tune$best.model
summary(best.linear)

##
## Call:
## best.svm(x = diabetes ~ ., data = dat[rowTrain, ], cost = exp(seq(-5,
##      2, len = 50)), kernel = "linear", scale = TRUE)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##      cost:  0.03741385
##
## Number of Support Vectors:  324
##
## ( 160 164 )
##
##
## Number of Classes:  2
##
## Levels:
##   pos neg
```

```

pred.linear <- predict(best.linear, newdata = dat[-rowTrain,])

confusionMatrix(data = pred.linear,
                 reference = dat$diabetes[-rowTrain])

```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction pos neg
##      pos   39   15
##      neg   28  110
##
##           Accuracy : 0.776
##           95% CI : (0.7104, 0.8329)
##      No Information Rate : 0.651
##      P-Value [Acc > NIR] : 0.0001183
##
##           Kappa : 0.4839
##
##  Mcnemar's Test P-Value : 0.0672525
##
##           Sensitivity : 0.5821
##           Specificity : 0.8800
##      Pos Pred Value : 0.7222
##      Neg Pred Value : 0.7971
##           Prevalence : 0.3490
##      Detection Rate : 0.2031
##      Detection Prevalence : 0.2812
##      Balanced Accuracy : 0.7310
##
##      'Positive' Class : pos
##


```

```

# we have 8 predictors but can only plot 2 at once
plot(best.linear, dat[rowTrain,],
     glucose ~ mass,
     slice = list(pregnant = 5, triceps = 20,
                  insulin = 20, pressure = 75,
                  pedigree = 1, age = 50),
     grid = 100) #higher grid the more defined line

```

## SVM classification plot



*#x obs are support vectors => critical in decide decision boundary*


### Radial kernel RBF

Support vector machines can construct classification boundaries that are nonlinear in shape. We use the radial kernel.

```
set.seed(1)
radial.tune <- tune.svm(diabetes ~ . ,
  data = dat[rowTrain,],
  kernel = "radial",
  cost = exp(seq(-1,4,len=10)),
  gamma = exp(seq(-6,-2,len=10)))

plot(radial.tune, transform.y = log, transform.x = log,
  color.palette = terrain.colors) #green means misclassification is lower
```

## Performance of 'svm'



```
# summary(radial.tune)
```

```
best.radial <- radial.tune$best.model
summary(best.radial)
```


```
##
## Call:
## best.svm(x = diabetes ~ ., data = dat[rowTrain, ], gamma = exp(seq(-6,
##      -2, len = 10)), cost = exp(seq(-1, 4, len = 10)), kernel = "radial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##      cost:  1.947734
##
## Number of Support Vectors:  327
##
## ( 164 163 )
##
##
## Number of Classes:  2
##
## Levels:
##   pos neg
```

```
pred.radial <- predict(best.radial, newdata = dat[-rowTrain,])
```

```
confusionMatrix(data = pred.radial,
  reference = dat$diabetes[-rowTrain])
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction pos neg
##      pos   39   16
##      neg   28  109
##
##           Accuracy : 0.7708
##           95% CI : (0.7048, 0.8283)
##      No Information Rate : 0.651
##      P-Value [Acc > NIR] : 0.0002216
##
##           Kappa : 0.4738
##
## Mcnemar's Test P-Value : 0.0972544
##
##           Sensitivity : 0.5821
##           Specificity : 0.8720
##           Pos Pred Value : 0.7091
##           Neg Pred Value : 0.7956
##           Prevalence : 0.3490
##           Detection Rate : 0.2031
##           Detection Prevalence : 0.2865
##           Balanced Accuracy : 0.7270
##
##           'Positive' Class : pos
##
#visualize decision boundary
plot(best.radial, dat[rowTrain,],
      glucose ~ mass,
      slice = list(pregnant = 5, triceps = 20,
                   insulin = 20, pressure = 75,
                   pedigree = 1, age = 50),
      grid = 100,
      symbolPalette = c("cyan", "darkblue"), #class labels
      color.palette = heat.colors) #background colors: rainbow, etc
```


## SVM classification plot



```
plot(best.radial, dat[rowTrain,],  
      glucose ~ mass,  
      slice = list(pregnant = 5, triceps = 20,  
                   insulin = 20, pressure = 75,  
                   pedigree = 1, age = 50),  
      grid = 200)
```



## SVM classification plot



## Using **kernlab** - another commonly used library in R for SVM

Check <https://cran.r-project.org/web/packages/kernlab/vignettes/kernlab.pdf> for more details. This has only one tuning parameter to show the syntax, not recommended for model training.

```
x_train <- as.matrix(dat[rowTrain, 1:8])
x_test  <- as.matrix(dat[-rowTrain, 1:8])

linear <- ksvm(x = x_train,
               y = dat$diabetes[rowTrain],
               type = "C-svc",
               kernel = "vanilladot",
               C = 1,
               scaled = TRUE)
```

```
## Setting default kernel parameters
```

```
pred.linear2 <- predict(linear, newdata = x_test)
```

```
# "?dots" for definition of kernel functions
```

```
set.seed(1)
rbf <- ksvm(x = x_train,
            y = dat$diabetes[rowTrain],
            type = "C-svc",
            kernel = "rbfdot",
            kpar = "automatic",
            C = 1)
```


```
pred.rbf <- predict(rbf, newdata = x_test)
```

## Using caret

```
ctrl <- trainControl(method = "cv")


# kernlab
set.seed(1)
svml.fit <- train(diabetes ~ . ,
                  data = dat[rowTrain,],
                  method = "svmLinear",
                  # preProcess = c("center", "scale"),
                  tuneGrid = data.frame(C = exp(seq(-2,5,len=20))),
                  trControl = ctrl)

plot(svml.fit, highlight = TRUE, xTrans = log)
```



```
# e1071
set.seed(1)
svml.fit2 <- train(diabetes ~ . ,
                  data = dat[rowTrain,],
                  method = "svmLinear2",
                  tuneGrid = data.frame(cost = exp(seq(-2,5,len=20))),
                  trControl = ctrl)


plot(svml.fit2, highlight = TRUE, xTrans = log)
```



```
svmr.grid <- expand.grid(C = exp(seq(-1,4,len=10)),
                        sigma = exp(seq(-8,0,len=10)))

# tunes over both cost and sigma
set.seed(1)
svmr.fit <- train(diabetes ~ . , dat,
                  subset = rowTrain,
                  method = "svmRadialSigma",
                  preProcess = c("center", "scale"),
                  tuneGrid = svmr.grid,
                  trControl = ctrl)

plot(svmr.fit, highlight = TRUE)
```



```
# tune over cost and uses a single value of sigma based on kernlab's sigest function
set.seed(1)
svmr.fit2 <- train(diabetes ~ . , dat,
  subset = rowTrain,
  method = "svmRadialCost",
  preprocess = c("center", "scale"),
  tuneGrid = data.frame(C = exp(seq(-4,2,len=10))),
  trControl = ctrl)

# Platt's probabilistic outputs; use with caution
set.seed(1)
svmr.fit3 <- train(diabetes ~ . , dat,
  subset = rowTrain,
  method = "svmRadialCost",
  preprocess = c("center", "scale"),
  tuneGrid = data.frame(C = exp(seq(-4,2,len=10))),
  trControl = ctrl,
  prob.model = TRUE)
# predict(svmr.fit3, newdata = x_test, type = "prob")


set.seed(1)
rpart.fit <- train(diabetes ~ . , dat,
  subset = rowTrain,
  method = "rpart",
  tuneLength = 50,
  trControl = ctrl)

resamp <- resamples(list(svmr = svmr.fit, svmr2 = svmr.fit2,
  svml = svml.fit, svml2 = svml.fit2,
```

```

                                rpart = rpart.fit))
bwplot(resamp)

```



We finally look at the test data performance.

```

pred.svml <- predict(svml.fit, newdata = dat[-rowTrain,])
pred.svmr <- predict(svmr.fit, newdata = dat[-rowTrain,])

```

```

confusionMatrix(data = pred.svml,
                 reference = dat$diabetes[-rowTrain])

```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction pos neg
##      pos  40  17
##      neg  27 108
##
##           Accuracy : 0.7708
##           95% CI : (0.7048, 0.8283)
##      No Information Rate : 0.651
##      P-Value [Acc > NIR] : 0.0002216
##
##           Kappa : 0.4776
##
##  McNemar's Test P-Value : 0.1748444
##
##           Sensitivity : 0.5970
##           Specificity : 0.8640
##      Pos Pred Value : 0.7018
##      Neg Pred Value : 0.8000

```

```
##           Prevalence : 0.3490
##           Detection Rate : 0.2083
##       Detection Prevalence : 0.2969
##           Balanced Accuracy : 0.7305
##
##       'Positive' Class : pos
##
confusionMatrix(data = pred.svmr,
                 reference = dat$diabetes[-rowTrain])

## Confusion Matrix and Statistics
##
##           Reference
## Prediction pos neg
##           pos  41  17
##           neg  26 108
##
##           Accuracy : 0.776
##           95% CI : (0.7104, 0.8329)
##       No Information Rate : 0.651
##       P-Value [Acc > NIR] : 0.0001183
##
##           Kappa : 0.4912
##
## Mcnemar's Test P-Value : 0.2224692
##
##           Sensitivity : 0.6119
##           Specificity : 0.8640
##       Pos Pred Value : 0.7069
##       Neg Pred Value : 0.8060
##           Prevalence : 0.3490
##       Detection Rate : 0.2135
##       Detection Prevalence : 0.3021
##       Balanced Accuracy : 0.7380
##
##       'Positive' Class : pos
##
```

## Understanding your models with DALEX

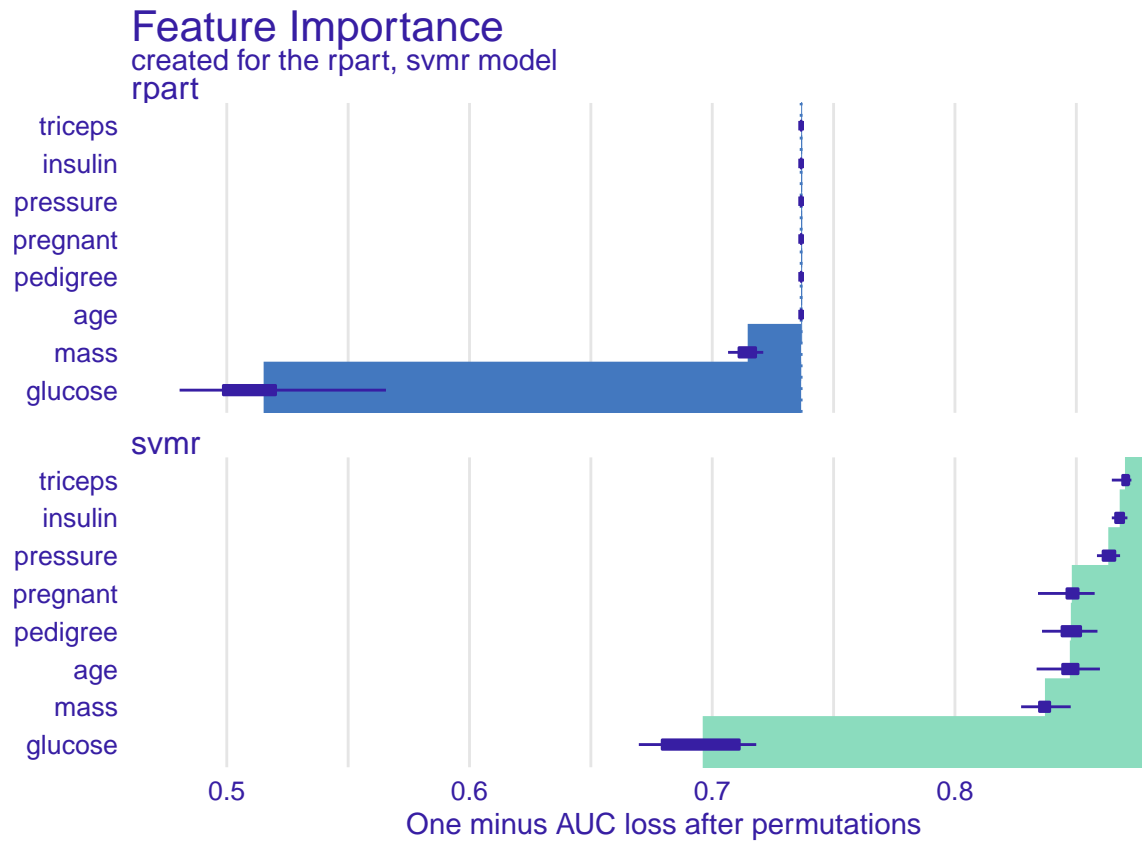
```
explainer_rpart <- explain(rpart.fit,
                          label = "rpart",
                          data = x_train,
                          y = as.numeric(dat$diabetes[rowTrain] == "pos"),
                          verbose = FALSE)
```

*# SVM does not output probabilities, this is using Platt's output and is just for illustration of DALEX*

```
explainer_svm <- explain(svmr.fit3,
                        label = "svmr",
                        data = x_train,
                        y = as.numeric(dat$diabetes[rowTrain] == "pos"),
                        verbose = FALSE)
```

```
# variable importance
vi_rpart <- model_parts(explainer_rpart)
vi_svm <- model_parts(explainer_svm)

plot(vi_rpart, vi_svm)
```



```
# PDP
pdp_svm <- model_profile(explainer_svm,
  variable = "glucose",
  type = "partial")
pdp_rpart <- model_profile(explainer_rpart,
  variable = "glucose",
  type = "partial")
plot(pdp_svm, pdp_rpart)
```

## Partial Dependence profile

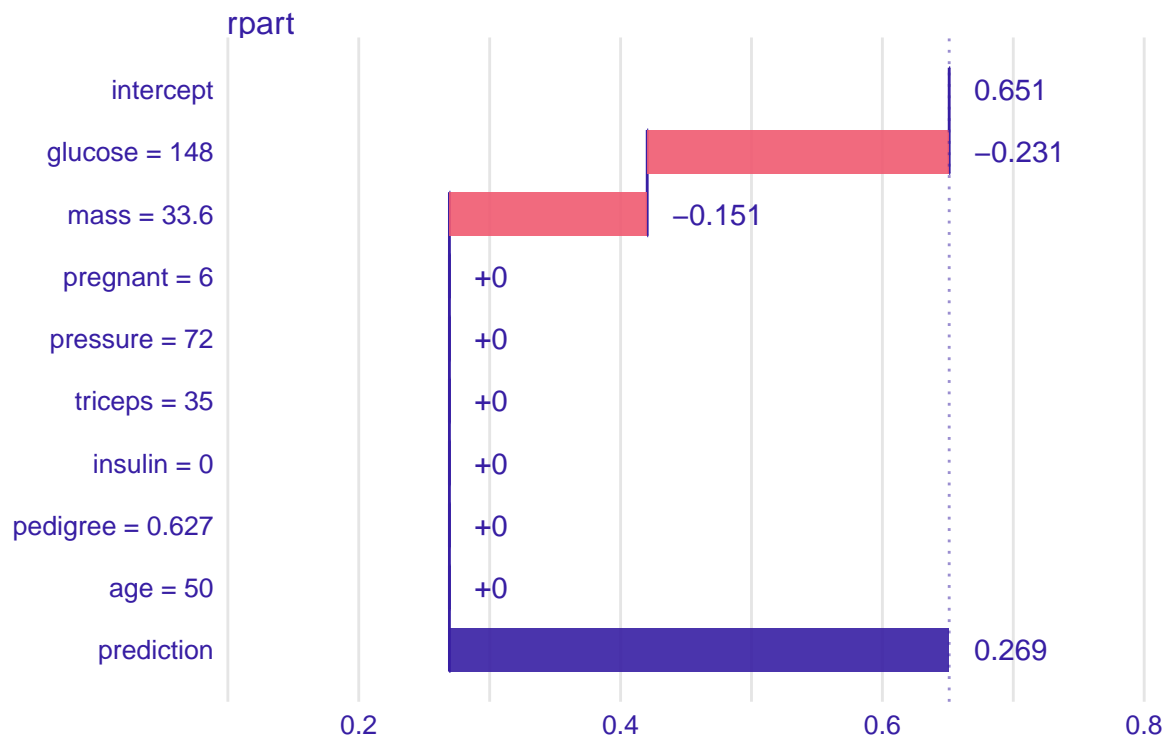
Created for the svmr, rpart model



```
# bread down
pb_svmr <- predict_parts(explainer_svmr,
                        new_observation = dat[1,],
                        type = "break_down")
pb_rpart <- predict_parts(explainer_rpart,
                        new_observation = dat[1,],
                        type = "break_down")
plot(pb_rpart)
```




## Break Down profile



```
plot(pb_svm)
```

## Break Down profile



## Regression

Predict a baseball player's salary on the basis of various statistics associated with performance in the previous year. Use `?Hitters` for more details.

```
data(Hitters)
Hitters <- na.omit(Hitters)

set.seed(2021)
trRows <- createDataPartition(Hitters$Salary,
                               p = .75,
                               list = F)
```

### eps-regression

```
set.seed(1)
svml.fit <- tune.svm(Salary ~ . ,
                    data = Hitters[trRows,],
                    kernel = "linear",
                    epsilon = exp(seq(-5,0,len=20)))

set.seed(1)
svmr.fit <- tune.svm(Salary ~ . ,
                    data = Hitters[trRows,],
                    kernel = "radial",
                    epsilon = exp(seq(-5,0,len=10)),
                    gamma = exp(seq(-6,-2,len=10)))

svmr.pred <- predict(svmr.fit$best.model, newdata = Hitters[-trRows,])
RMSE(svmr.pred, Hitters$Salary[-trRows])

## [1] 292.2582
```