

# Clustering and PCA

Yifei Sun

## Contents

<b>K means clustering</b>	<b>3</b>
<b>Hierarchical clustering</b>	<b>4</b>
<b>PCA</b>	<b>6</b>
PCA - Biplot . . . . .	8
<b>PCA for image compression</b>	<b>10</b>

```
library(ISLR)
library(factoextra) #visualizing cluster and PCA
library(gridExtra) #arrange multiple plots
library(corrplot)
library(RColorBrewer) #to generate colors
library(gplots)
library(jpeg)
library(tidyverse)
```

The dataset we use contains data on 166 first generation Pokemon, including their names and basic stats: HP, Attack, Defense, Special Attack, Special Defense, and Speed. The data is from Kaggle (<https://www.kaggle.com/abcsds/pokemon>). We will apply unsupervised learning methods on this data. The list of Pokemon can be found at (<https://pokemondb.net/pokedex/national>).

```
dat <- read.csv("data/Pokemon.csv")
head(dat)
```

```
##           Name HitPoints Attack Defense SpecialAttack SpecialDefense
## 1      Bulbasaur      45     49     49             65             65
## 2      Ivysaur      60     62     63             80             80
## 3      Venusaur      80     82     83            100            100
## 4 VenusaurMega Venusaur      80    100    123            122            120
## 5      Charmander      39     52     43             60             50
## 6      Charmeleon      58     64     58             80             65
##   Speed Legendary
## 1     45      FALSE
## 2     60      FALSE
## 3     80      FALSE
## 4     80      FALSE
## 5     65      FALSE
## 6     80      FALSE
```

```
dim(dat)
```

```
## [1] 166   8
```

```
dat1 <- dat[,2:7]
head(dat1)
```

```
##   HitPoints Attack Defense SpecialAttack SpecialDefense Speed
## 1      45     49     49             65             65    45
## 2      60     62     63             80             80    60
## 3      80     82     83            100            100    80
## 4      80    100    123            122            120    80
## 5      39     52     43             60             50    65
## 6      58     64     58             80             65    80
```

```
dat1 <- scale(dat1)
head(dat1)
```

```
##           HitPoints      Attack      Defense SpecialAttack SpecialDefense      Speed
## [1,] -0.7394793 -0.8989689 -0.7632830   -0.1980097   -0.1603732 -0.9295212
## [2,] -0.2066947 -0.4761322 -0.2744790    0.2375417    0.4277405 -0.4240598
## [3,]  0.5036847  0.1743859  0.4238124    0.8182769    1.2118920  0.2498887
## [4,]  0.5036847  0.7598521  1.8203953    1.4570855    1.9960435  0.2498887
## [5,] -0.9525931 -0.8013912 -0.9727705   -0.3431935   -0.7484868 -0.2555726
## [6,] -0.2777327 -0.4110804 -0.4490519    0.2375417   -0.1603732  0.2498887
```


```
rownames(dat1) <- dat[,1]
head(dat1)
```

```
##              HitPoints      Attack      Defense SpecialAttack
## Bulbasaur      -0.7394793 -0.8989689 -0.7632830    -0.1980097
## Ivysaur        -0.2066947 -0.4761322 -0.2744790     0.2375417
## Venusaur       0.5036847  0.1743859  0.4238124     0.8182769
## VenusaurMega Venusaur 0.5036847  0.7598521  1.8203953     1.4570855
## Charmander    -0.9525931 -0.8013912 -0.9727705    -0.3431935
## Charmeleon    -0.2777327 -0.4110804 -0.4490519     0.2375417
##              SpecialDefense      Speed
## Bulbasaur      -0.1603732 -0.9295212
## Ivysaur         0.4277405 -0.4240598
## Venusaur       1.2118920  0.2498887
## VenusaurMega Venusaur 1.9960435  0.2498887
## Charmander    -0.7484868 -0.2555726
## Charmeleon    -0.1603732  0.2498887
```

## K means clustering

Partitioning methods such as k-means clustering require the users to specify the number of clusters to be generated. The function `fviz_nbclust()` determines and visualizes the optimal number of clusters using different methods: within cluster sums of squares, average silhouette and gap statistics. We use average silhouette, and the greater the silhouette value the better.

```
fviz_nbclust(dat1,
  FUNcluster = kmeans,
  method = "silhouette") #optimal number = 2
```



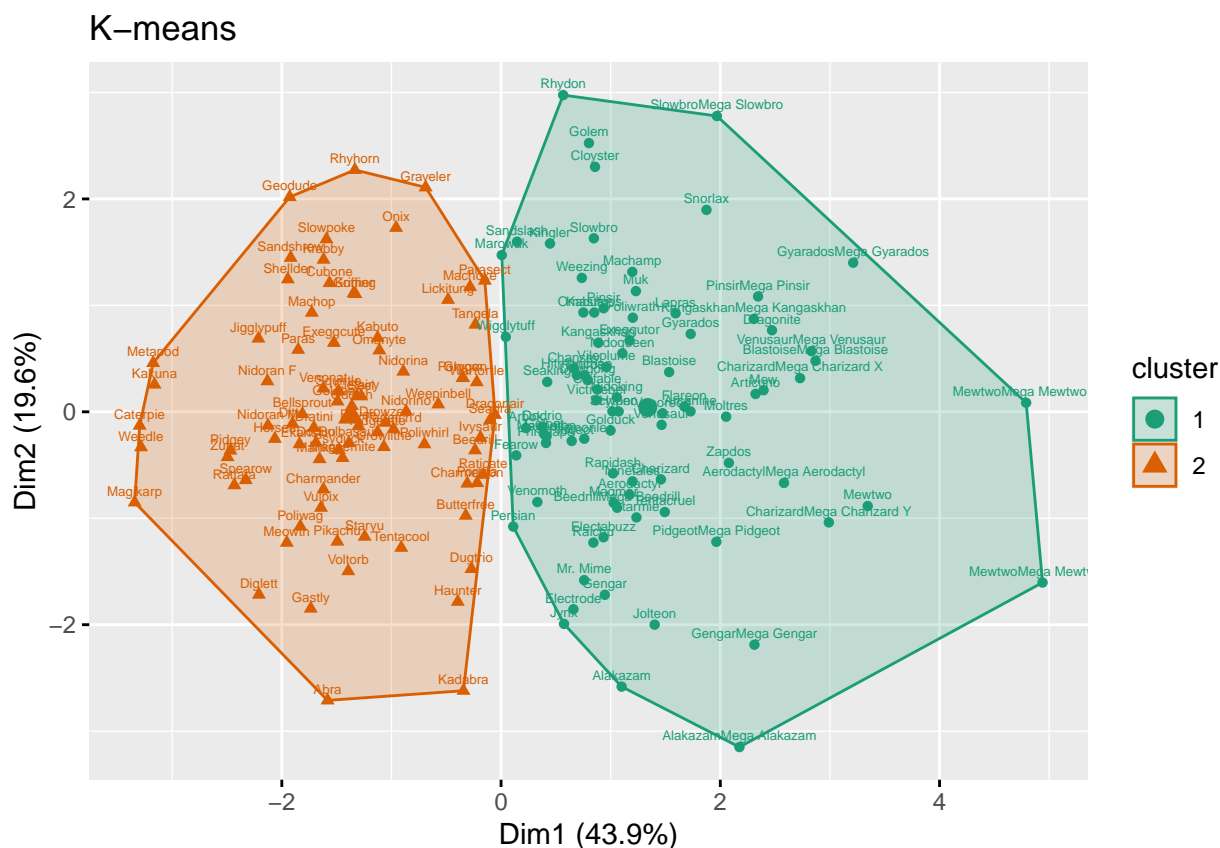
```
set.seed(1)
km <- kmeans(dat1, centers = 2, nstart = 20) #baseR function
```

The function `fviz_cluster()` provides ggplot2-based visualization of partitioning methods including K

means. Observations are represented by points in the plot, using principal components if  $p > 2$ . An ellipse is drawn around each cluster.

```
km_vis <- fviz_cluster(list(data = dat1, cluster = km$cluster),
  ellipse.type = "convex",
  geom = c("point", "text"),
  labels = 5,
  palette = "Dark2") + labs(title = "K-means")
```

km\_vis



## Hierarchical clustering

We can also apply hierarchical clustering on this data. Here we use the Euclidean distance and different types of linkage.

```
hc.complete <- hclust(dist(dat1), method = "complete") #default setting
hc.average <- hclust(dist(dat1), method = "average")
hc.single <- hclust(dist(dat1), method = "single")
hc.centroid <- hclust(dist(dat1), method = "centroid")
```


*# average and centroid are usually between single and complete linkage*

The function `fviz_dend()` can be applied to visualize the dendrogram.

```
fviz_dend(hc.complete, k = 4,
  cex = 0.3,
  palette = "jco",
```

```
color_labels_by_k = TRUE,
rect = TRUE, rect_fill = TRUE, rect_border = "jco",
labels_track_height = 2.5)
```

## Cluster Dendrogram



```
ind4.complete <- cutree(hc.complete, 4)
```


```
# Who are in the fourth cluster?
dim(dat[ind4.complete == 4,])
```

```
## [1] 6 8
```

To display more details, we show the heatmap of the data.

```
#display.brewer.all(n=NULL, type="all", select=NULL, exact.n=TRUE)
col1 <- colorRampPalette(brewer.pal(9, "GnBu"))(100)
col2 <- colorRampPalette(brewer.pal(3, "Spectral"))(2)

heatmap.2(t(dat1),
  col = col1, keysize=.8, key.par = list(cex=.5),
  trace = "none", key = TRUE, cexCol = 0.75,
  labCol = as.character(dat[,1]),
  ColSideColors = col2[as.numeric(dat[, "Legendary"])+1],
  margins = c(10, 10))
```



## PCA

The function `prcomp()` can be used to perform PCA. PCA finds low-dimensional features to represent data ( $p < n$ )

```
pca <- prcomp(dat1)
pca$rotation
```

```
##          PC1      PC2      PC3      PC4      PC5
## HitPoints  0.3638022  0.2862972 -0.72425610  0.078135517 -0.42663445
## Attack    0.4363956  0.3149994  0.27779372  0.573757403 -0.12333177
## Defense   0.3031184  0.5812622  0.48929801 -0.361032760  0.03453292
## SpecialAttack 0.4378985 -0.3119077  0.03743076 -0.654754892 -0.29724064
## SpecialDefense 0.5204254 -0.1331800 -0.25037584  0.006718302  0.80453025
## Speed     0.3503261 -0.6049135  0.30786700  0.324966961 -0.25682461
##          PC6
## HitPoints  -0.27020899
## Attack     0.53736079
## Defense   -0.44643458
## SpecialAttack 0.43874986
## SpecialDefense -0.03766086
## Speed     -0.49498168
```

```
# pca$x returns z scores
```

```
pca$sdev #sd(Zk) = d_k/sqrt(n-1), in decreasing order
```

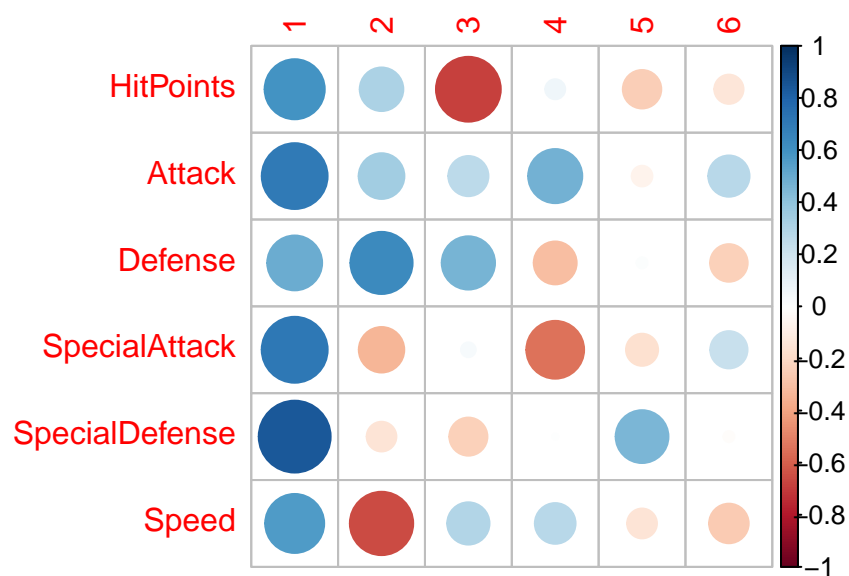
```
## [1] 1.6238460 1.0848056 0.9487926 0.8345883 0.5670204 0.5177487
```

```
pca$rotation %>% diag(pca$sdev) #correlation loading: between X (original p features) and Z (transform
```

```
##          [,1]      [,2]      [,3]      [,4]      [,5]
```

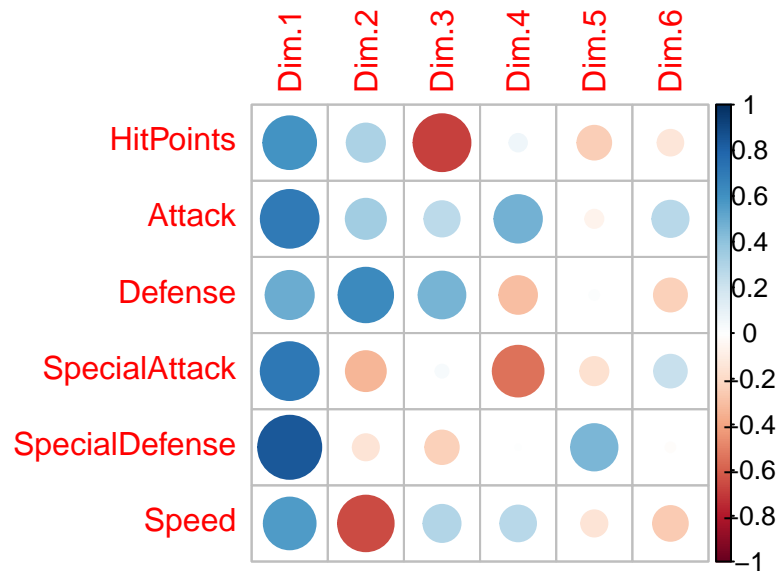
```
## HitPoints      0.5907587  0.3105768 -0.68716880  0.065210990 -0.24191044
## Attack        0.7086393  0.3417131  0.26356862  0.478851234 -0.06993163
## Defense       0.4922176  0.6305564  0.46424231 -0.301313729  0.01958087
## SpecialAttack 0.7110798 -0.3383592  0.03551403 -0.546450792 -0.16854151
## SpecialDefense 0.8450908 -0.1444744 -0.23755473  0.005607016  0.45618508
## Speed        0.5688756 -0.6562135  0.29210193  0.271213634 -0.14562479
##              [,6]
## HitPoints    -0.13990035
## Attack       0.27821784
## Defense     -0.23114091
## SpecialAttack 0.22716216
## SpecialDefense -0.01949886
## Speed       -0.25627611
```

```
corrplot(pca$rotation %*% diag(pca$sdev))
```




```
var <- get_pca_var(pca)
```

```
corrplot(var$cor)
```



The function `fviz_eig()` plots the eigenvalues/variances against the number of dimensions.

```
# Scree plot: p = 6. Want to find "elbow": where after this pc the line flatten out
fviz_eig(pca, addlabels = TRUE)
```




## PCA - Biplot

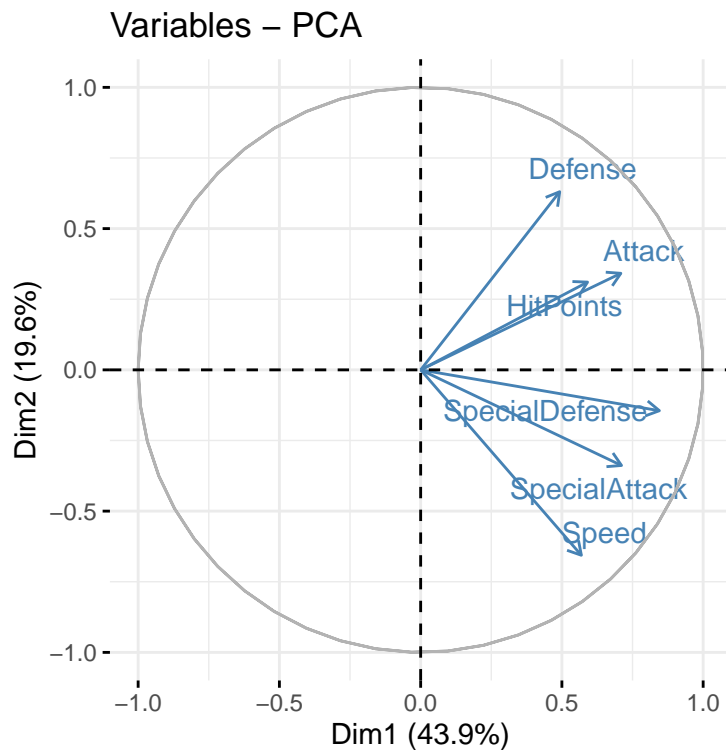
The function `fviz_pca_biplot()` can be used to obtain the biplot of individuals and variables.

```
fviz_pca_biplot(pca, axes = c(1,2),
  habillage = ifelse(dat$Legendary==TRUE, "Legendary", "Not legendary"),
  label = c("var"),
  addEllipses = TRUE)
```






```
fviz_pca_var(pca, col.var = "steelblue", repel = TRUE)
```



```
fviz_pca_ind(pca,
  habillage = ifelse(dat$Legendary==TRUE,"Legendary","Not legendary"),
  label = "none",
  addEllipses = TRUE)
```




## PCA for image compression

```
img <- readJPEG("image.jpeg")
dim(img)
```

```
## [1] 948 1685 3
```

```
knitr::include_graphics("image.jpeg")
```



```

r <- img[,1]
g <- img[,2]
b <- img[,3]

img.r.pca <- prcomp(r, center = FALSE)
img.g.pca <- prcomp(g, center = FALSE)
img.b.pca <- prcomp(b, center = FALSE)

rgb.pca <- list(img.r.pca, img.g.pca, img.b.pca)

# Approximate X with XV_kV_k^T
compress <- function(pr, k)
{
  compressed.img <- pr$x[,1:k] %*% t(pr$rotation[,1:k])
  compressed.img
}

# Using first 20 PCs
pca20 <- sapply(rgb.pca,
                compress,
                k = 20,
                simplify = "array") #can use map function

writeJPEG(pca20, "pca20.jpeg")
knitr::include_graphics("pca20.jpeg")

```

